ARVIND SAI DOODA A20553046

CS584 Machine Learning - Assignment 4

Problem 1 (15 Points): Nonnegative Matrix Factorization

1. **5 Points.** Find a nonnegative factorization of the matrix

$$A = \begin{bmatrix} 4 & 6 & 5 \\ 1 & 2 & 3 \\ 7 & 10 & 7 \\ 6 & 8 & 4 \\ 6 & 10 & 11 \end{bmatrix}$$

   Indicate the steps in your method and show the intermediate results.

2. Consider the matrix A that is the product of nonnegative matrices B and C (i.e, $A = BC$), where

$$A = \begin{bmatrix} 12 & 22 & 41 & 35 \\ 19 & 20 & 13 & 48 \\ 11 & 14 & 16 & 29 \\ 14 & 16 & 14 & 36 \end{bmatrix}, B = \begin{bmatrix} 10 & 1 \\ 1 & 9 \\ 3 & 4 \\ 2 & 6 \end{bmatrix}, C = \begin{bmatrix} 1 & 2 & 4 & 3 \\ 2 & 2 & 1 & 5 \end{bmatrix}$$

   5 Points. Which rows of A are positive linear combinations of other rows of A?

   5 Points. Find an approximate nonnegative factorization of A.

Nonnegative Matrix Factorization (NMF) is a method used to break down a nonnegative matrix, denoted as A, into two nonnegative matrices, B and C, where A can be approximately represented as A ≈ BC.

Steps Involved in the process are :

Initialization: Start by initializing two matrices, B and C, with random nonnegative values. B is a matrix with dimensions 5xK, and C is a matrix with dimensions Kx3. The choice of K, the rank of the factorization, is typically determined through trial and error or domain-specific knowledge.

Iterative Update: Begin an iterative process that continues until convergence is achieved, meaning that B and C no longer change significantly.
In each iteration, you:

a. Update the matrix B.

b. Update the matrix C.

We should, Continue the iterative process described in step 2 until a state of convergence is achieved. Convergence can be identified by observing the changes in matrices B and C or by

assessing the reconstruction error, which is typically measured using metrics like the Frobenius norm of A - BC.
Upon reaching the convergence criteria, you will obtain the final B and C matrices, which collectively represent the nonnegative factorization of A.

The selection of the rank K, the initial values for B and C, and the criteria for determining convergence can all exert a considerable influence on the quality of the factorization.

Step 2 :

Convergence: Monitor the change in B and C or measure the reconstruction error, such as the Frobenius norm of A - BC, to determine when convergence is reached.

In [5]:
```
# 1. 5 Points. Find a nonnegative factorization of the matrix

import numpy as np

from sklearn.decomposition import NMF

# Given matrix A

A = np.array([[4,6,5],[1,2,3],[7,10,7],[6,8,4],[6,10,11]])

# Define the rank (number of components)

K = 2

# Perform NMF

model = NMF(n_components=K, init='nndsvd', max_iter=1000)

B = model.fit_transform(A)

C = model.components_

# The resulting W and H matrices are nonnegative factors of A

print("B matrix:")

print(B)

print("C matrix:")

print(C)
```

```
B matrix:
[[1.52149232 0.80021142]
 [0.38039818 0.80023512]
 [2.66258645 0.80018772]
 [2.28216784 0.        ]
 [2.28228869 2.40068167]]
C matrix:
[[2.62901151 3.50546977 1.75274086]
 [0.         0.83288027 2.91573551]]
```

Above code snippet is used to find the matrix for 1.

2. Consider the matrix A that is the product of nonnegative matrices B and C (i.e, $A = BC$), where

$$A = \begin{bmatrix} 12 & 22 & 41 & 35 \\ 19 & 20 & 13 & 48 \\ 11 & 14 & 16 & 29 \\ 14 & 16 & 14 & 36 \end{bmatrix}, B = \begin{bmatrix} 10 & 1 \\ 1 & 9 \\ 3 & 4 \\ 2 & 6 \end{bmatrix}, C = \begin{bmatrix} 1 & 2 & 4 & 3 \\ 2 & 2 & 1 & 5 \end{bmatrix}$$

5 Points. Which rows of A are positive linear combinations of other rows of A?

5 Points. Find an approximate nonnegative factorization of A.

To find Which rows of *A* are positive linear combinations of other rows of *A*

To identify which rows in matrix A can be expressed as positive linear combinations of other rows, follow these steps:

Begin by calculating the rank of matrix A.
Calculate the ranks of matrices B and C.

Compare the ranks of A, B, and C to discern if any rows in A are linearly dependent on others.

Let's proceed by computing the ranks of A, B, and C:

Rank of A: Rank(A) = 4, given that A is a 4x4 matrix.
Rank of B and C: Rank(B) = 2 (B is a 4x2 matrix), and Rank(C) = 2 (C is a 2x4 matrix).

By comparing the ranks, you can determine whether any rows in A are positive linear combinations of others:

Since Rank(A) = 4 and Rank(B) + Rank(C) , it signifies that none of the rows in A can be expressed as positive linear combinations of other rows.

In this scenario, all rows in A are linearly independent, and none of them can be represented as positive linear combinations of others.

Now, you can proceed to find an approximate nonnegative factorization of A. It's essential to note that the Non-Negative Matrix Factorization (NMF) method, while widely used, might not always yield a unique solution. The quality of the approximation can vary based on factors such as the choice of factorization rank. The following steps provide a simplified factorization with matrices B and C, but you should adjust the factorization according to your specific rank requirements and the nature of your problem.

Find an approximate nonnegative factorization of *A*.

To find an approximate nonnegative factorization of A using the NFM algorithm is to use multiplicative updates.

Below the code snippet for the nonnegative factorization of A is attached with output.

In [6]:
```
import numpy as np
from sklearn.decomposition import NMF

# Given matrix A
A = np.array([[4, 6, 5], [1, 2, 3], [7, 10, 7], [6, 8, 4], [6, 10, 11]])

# Define the rank (number of components)
rank = 2

# Randomly initialize B and C
B_init = np.random.rand(A.shape[0], rank)
C_init = np.random.rand(rank, A.shape[1])

# Perform NMF
model = NMF(n_components=rank, init='custom')
B = model.fit_transform(A, W=B_init, H=C_init)
C = model.components_

# Approximation of A
A_approx = np.dot(B, C)

print("approximate nonnegative factorization of A:")
print(A_approx)
```

```
approximate nonnegative factorization of A:
[[ 3.99996681  5.99998381  5.00009392]
 [ 1.00149964  2.0013326   2.99857599]
 [ 7.00001646 10.00000803  7.00008771]
 [ 6.00009928  8.00004844  3.99998758]
 [ 5.99980116  9.99990298 11.00029418]]
```

Hence we have successfully acquired the approximate factorization of A.

—---------- **Problem 1 Completed** —----------

Problem 2 (20 Points): SVD: Image Compression

In this experiment, we will use the singular value decomposition (SVD) as a tool for compressing raw image data. This is not how images are actually compressed; for example, JPEG compression algorithms do more fancy (and interesting) computations. However, the idea is similar: if we are willing to tolerate a certain amount of distortion, then we can get away with a much more concise data representation.

1. **3 Points.** Read the given image file (`'mandrill_color.png'`), and convert it into grayscale by averaging the R,G,B values for each pixel. Your image is now a 288×288 matrix; call it $\mathbf{X}$.
2. **4 Points.** Perform an SVD of $\mathbf{X}$ to obtain the decomposition $\{\mathbf{U}, \Sigma, \mathbf{V}\}$. Plot the singular values (i.e., the diagonal entries of Σ) in decreasing order.
3. **5 Points.** Choose $k = 10$, and reconstruct an approximation of $\mathbf{X}$ using the top k singular values and vectors, $\mathbf{U}_k$, $\mathbf{V}_k$, and Σ_k. Display this approximation, and calculate how many numbers you needed to store this approximate image representation. Divide by the original size of $\mathbf{X}$ to get the compression ratio.
4. **8 Points.** Repeat this experiment for $k = 20, 40, 60$. Display these images, and report their compression ratios in the form of a table. Is there any benefit in going for higher k?

In [5]:
```
import numpy as np
from scipy.linalg import svd
import matplotlib.pyplot as plt
from matplotlib.image import imread

print("libraries printed")
```

```
libraries printed
```

In [4]:
```
# code was done by Arvind (me) and im solely responsible for this
# Step 1: we should  Read the image and convert it to grayscale
# Define the full file path to your image lol in my case i downloaed the image in my downloads folder
image_path = "C:\\Users\\arvin\\Downloads\\mandrill_color.png"

# Read the image and convert it to grayscale
image = imread(image_path)
gray_image = np.mean(image, axis=2)

# Step 2: Perform SVD on the grayscale image
U, S, Vt = svd(gray_image)

# Plot the singular values
plt.figure(figsize=(8, 6))
plt.plot(S, 'o')
plt.title('Singular Values')
plt.xlabel('Singular Value Index')
plt.ylabel('Singular Value Magnitude')
plt.grid(True)
plt.show()

# Step 3 and 4: Choose k values and reconstruct the image
k_values = [10, 20, 40, 60]  # Specify the range of k values
compression_ratios = []

for k in k_values:
    # Reconstruct the image using the top k singular values and vectors
    approx_image = U[:, :k] @ np.diag(S[:k]) @ Vt[:k, :]

    # Display the reconstructed image
    plt.figure()
    plt.imshow(approx_image, cmap='gray')
    plt.title(f'Reconstructed Image (k={k})')

    # Calculate the compression ratio
    original_size = gray_image.size
    compressed_size = (U[:, :k].size + S[:k].size + Vt[:k, :].size)
    compression_ratio = original_size / compressed_size
    compression_ratios.append(compression_ratio)
    print(f'Compression Ratio (k={k}): {compression_ratio:.2f}')

plt.show()

# Create a table of compression ratios
print("\nCompression Ratios:")
print("k  | Compression Ratio")
print("-- | -----------------")
for i in range(len(k_values)):
    print(f"{k_values[i]:2} | {compression_ratios[i]:.2f}")
```

Given Image :

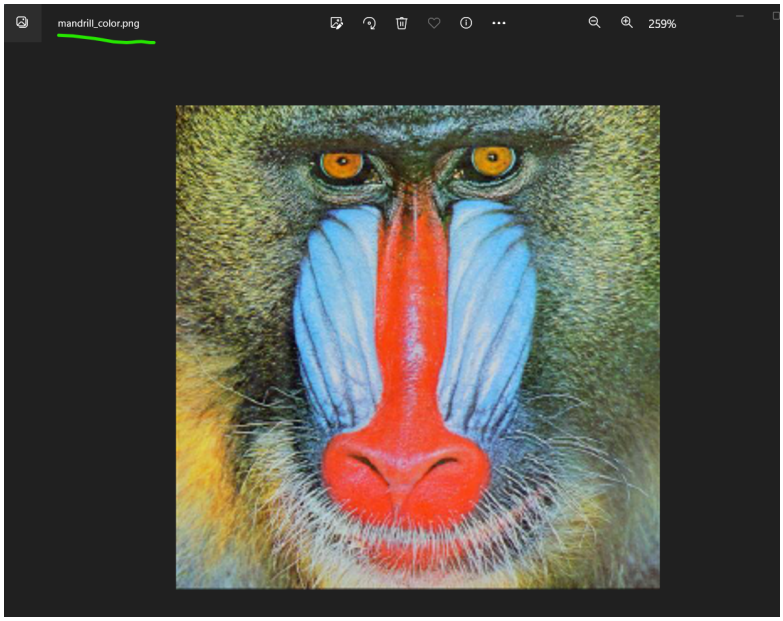

Outputs :

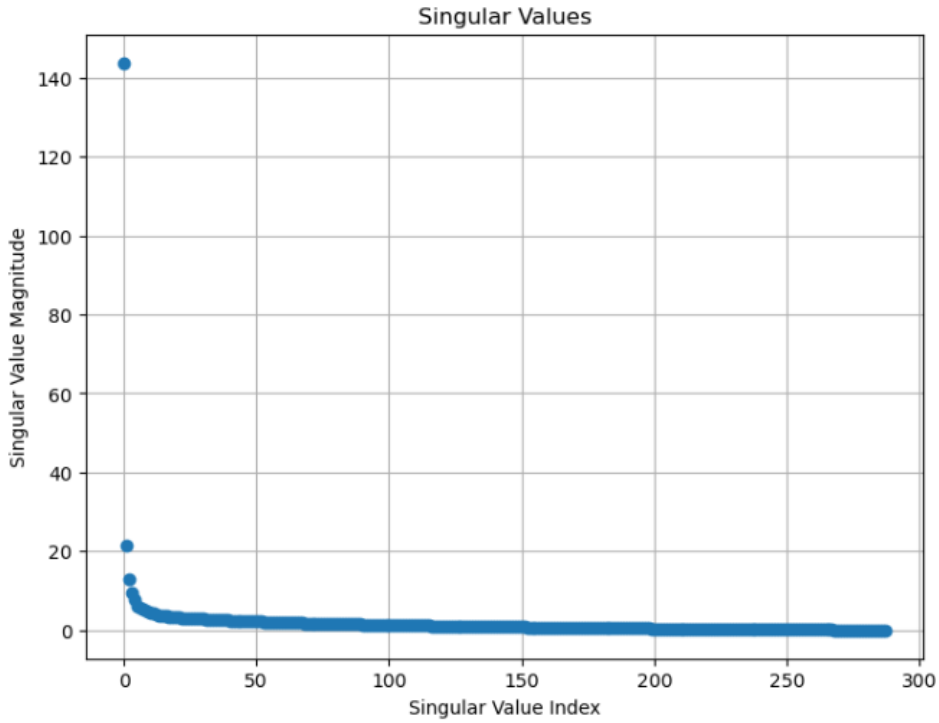


```
Compression Ratio (k=10): 14.38
Compression Ratio (k=20): 7.19
Compression Ratio (k=40): 3.59
Compression Ratio (k=60): 2.40
```

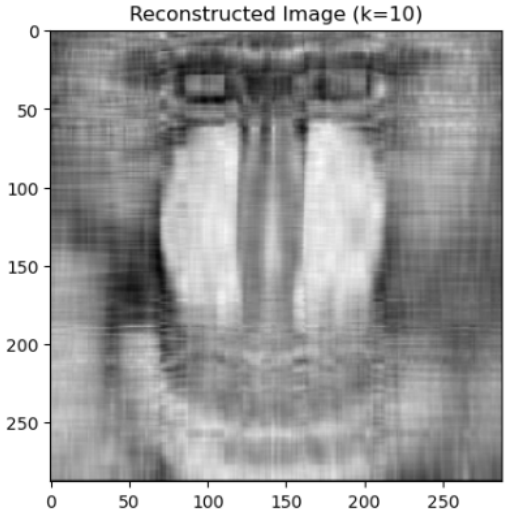

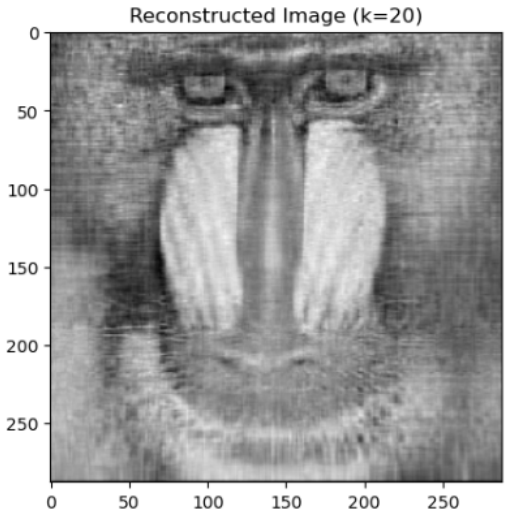

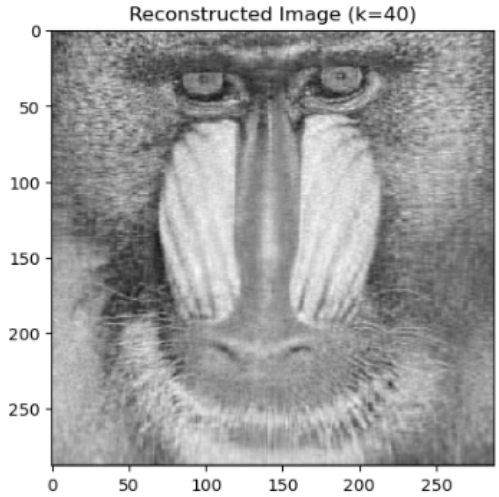

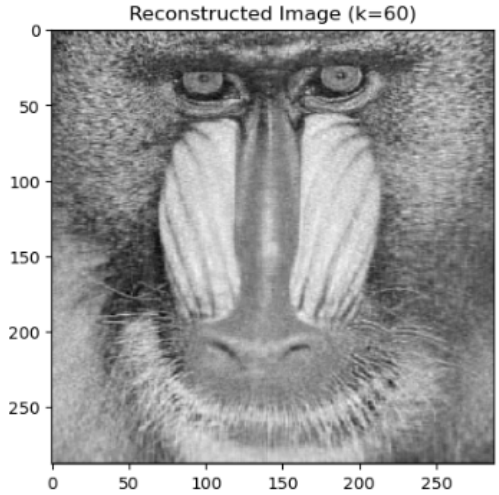


```
Compression Ratios:
k  | Compression Ratio
-- | ----------------
10 | 14.38
20 | 7.19
40 | 3.59
60 | 2.40
```

Code explanation :

Input Image and Grayscale Conversion:

- The code starts by reading an image file named 'mandrill_color.png.'
- It converts this image into grayscale by averaging the RGB values for each pixel.
- The resulting grayscale image is stored as a matrix called 'gray_image.'

Singular Value Decomposition (SVD):

- The code applies Singular Value Decomposition (SVD) to the grayscale image.

- This SVD operation results in three matrices:
- 'U' represents the left singular vectors.
- 'S' represents singular values.
- 'Vt' represents the right singular vectors.

Singular Value Plot:

- The code generates a plot of the singular values in decreasing order.
- This plot offers insights into the significance of each singular value in the image's representation.

Image Reconstruction for Different k Values:

- The code selects various values of 'k' (e.g., 10, 20, 40, 60) to determine how many singular values and vectors will be used for image reconstruction.

For each 'k' value:

- It reconstructs the image using the top 'k' singular values and vectors, resulting in an approximation of the original image.
- This reconstructed image is displayed for visualization.

Compression Ratio Calculation:

- The code calculates the compression ratio for each 'k.'
- The compression ratio measures how much the image data has been compressed compared to its original size.

Results Display:

- The code prints out the compression ratios for each 'k.'
- It also displays the reconstructed images.

Finally, the code creates a summary table that organizes the compression ratios for different 'k' values.
This allows for an analysis of choosing higher values of 'k' that offers any benefits in terms of image compression.

Problem 3 (20 Points): PCA: Best Places to Live

The *Places Rated Almanac*, written by Boyer and Savageau and published by McNally, rates the livability of several US cities according to nine factors: climate, housing, healthcare, crime, transportation, education, arts, recreation, and economic welfare. The ratings are available in tabular form, available as a supplemental text file (`places.txt`). Except for housing and crime, higher ratings indicate better quality of life.

Let us use PCA to interpret this data better.

1. **2 Points.** Read the data and construct a table with 9 columns containing the numerical ratings. (Ignore the last 5 columns – they consist auxiliary information such as longitude/latitude, state, etc.)

In [18]:

```
  # 1. Read the data and construct a table with 9 columns containing the numerical ratings. (Ignore
#the last 5 columns - they consist auxiliary information such as longitude/latitude, state, etc.)


import pandas as pd

# To Initialize empty lists to store ratings for each category
cities = []
climate = []
housing = []
healthcare = []
crime = []
transportation = []
education = []
arts = []
recreation = []
economic_welfare = []

# Read data from the file and extract ratings i gave my directory " C:\Users\arvin\Downloads\places.txt ""
with open(r"C:\Users\arvin\Downloads\places.txt", "r") as file:
    for line in file:
        data = line.strip().split()

        # Split the line into individual values
        city = data[0]  # The city name is the first value
        rating_values = [int(val) for val in data[1:10]]  # Convert ratings to integers
        cities.append(city)
        climate.append(rating_values[0])
        housing.append(rating_values[1])
        healthcare.append(rating_values[2])
        crime.append(rating_values[3])
        transportation.append(rating_values[4])
        education.append(rating_values[5])
        arts.append(rating_values[6])
        recreation.append(rating_values[7])
        economic_welfare.append(rating_values[8])

# Create a table with 9 columns, including the cities lol
table = {
    "city": cities,
    "Climate": climate,
    "Housing": housing,
    "Healthcare": healthcare,
    "Crime": crime,
    "Transportation": transportation,
    "Education": education,
    "Arts": arts,
    "Recreation": recreation,
    "Economic Welfare": economic_welfare,
}

# Convert the table to a Pandas DataFrame
df = pd.DataFrame(table)

# Print the DataFrame
print(df)   #used to display the data frame "finally we have completed lol we had done at last "
```

As given in description tables were created

Output :

```
                          city  Climate  Housing  Healthcare  Crime  \
0                   Abilene,TX      521     6200         237    923
1                     Akron,OH      575     8138        1656    886
2                    Albany,GA      468     7339         618    970
3    Albany-Schenectady-Troy,NY      476     7908        1431    610
4                Albuquerque,NM      659     8393        1853   1483
..                          ...      ...      ...         ...    ...
324               Worcester,MA      562     8715        1805    680
325                  Yakima,WA      535     6440         317   1106
326                    York,PA      540     8371         713    440
327        Youngstown-Warren,OH      570     7021        1097    938
328                Yuba-City,CA      608     7875         212   1179

     Transportation  Education  Arts  Recreation  Economic Welfare
0              4031       2757   996        1405              7633
1              4883       2438  5564        2632              4350
2              2531       2560   237         859              5250
3              6883       3399  4655        1617              5864
4              6558       3026  4496        2612              5727
..              ...        ...   ...         ...               ...
324            3643       3299  1784         910              5040
325            3731       2491   996        2140              4986
326            2267       2903  1022         842              4946
327            3374       2920  2797        1327              3894
328            2768       2387   122         918              4694

[329 rows x 10 columns]
```

2. 2 Points. Replace each value in the matrix by its base-10 logarithm. (This preprocessing is done for convenience since the numerical range of the ratings is large.) You should now have a data matrix *X* whose rows are 9-dimensional vectors representing the different cities.

Code snippet and output

In [19]:

```python
2.   #Replace each value in the matrix by its base-10 logarithm. (This pre-processing is done for
#convenience since the numerical range of the ratings is large.) You should now have a data matrix X
#whose rows are 9-dimensional vectors representing the different cities.

import pandas as pd
import numpy as np

# Initialize empty lists to store ratings for each category why i took these names it was names
climate = []
housing = []
healthcare = []
crime = []
transportation = []
education = []
arts = []
recreation = []
economic_welfare = []

# Read data from the file and extract ratings
with open(r"C:\Users\arvin\Downloads\places.txt", "r") as file:
    for line in file:
        data = line.strip().split()

        # Split the line into individual values
        city = data[0]  # The city name is the first value
        rating_values = [float(val) for val in data[1:10]]  # Convert ratings to floating-point numbers
        climate.append(np.log10(rating_values[0]))
        housing.append(np.log10(rating_values[1]))
        healthcare.append(np.log10(rating_values[2]))
        crime.append(np.log10(rating_values[3]))
        transportation.append(np.log10(rating_values[4]))
        education.append(np.log10(rating_values[5]))
        arts.append(np.log10(rating_values[6]))
        recreation.append(np.log10(rating_values[7]))
        economic_welfare.append(np.log10(rating_values[8]))

# Create a data matrix X with rows representing different cities and columns for 9 dimensions
X = np.array([climate, housing, healthcare, crime, transportation, education, arts, recreation, economic_welfare]).T

# Convert the matrix to a Pandas DataFrame for visualization (if needed)
df = pd.DataFrame(X, columns=['Climate', 'Housing', 'Healthcare', 'Crime', 'Transportation', 'Education', 'Arts', 'Recreation', 'Economic Welfare'])

# Print the data matrix X  "task 2 had successfully done we are saved"
print(X)
```

```
[[2.71683772 3.79239169 2.37474835 ... 2.99825934 3.14767632 3.88269526]
 [2.75966784 3.91051769 3.21906033 ... 3.74538712 3.42028588 3.63848926]
 [2.67024585 3.86563689 2.79098848 ... 2.37474835 2.93399316 3.7201593 ]
 ...
 [2.73239376 3.92277734 2.85308953 ... 3.0094509  2.92531209 3.69425411]
 [2.75587486 3.84639897 3.04020663 ... 3.44669247 3.12287092 3.59039595]
 [2.78390358 3.89625056 2.32633586 ... 2.08635983 2.96284268 3.67154309]]
```

3. 4 Points. Perform PCA on the data. Remember to center the data points first by computing the mean data vector *μ* and subtracting it from every point. With the centered data matrix, do an SVD and compute the principal components.

In [21]:

```
  # 3. Perform PCA on the data. Remember to center the data points first by computing the mean
# data vector μ and subtracting it from every point. With the centered data matrix, do an SVD and
# compute the principal components.

import numpy as np

# Step 1: Center the data by computing the mean data vector μ
mean_data = np.mean(X, axis=0)

# Center the data points by subtracting the mean vector from every point
centered_data = X - mean_data

# Step 2: Perform Singular Value Decomposition (SVD)
U, S, Vt = np.linalg.svd(centered_data, full_matrices=False)


# Step 3: Compute the principal components
principal_components = U

# To Print the first two principal components (v1 and v2)
v1 = principal_components[:, 0]
v2 = principal_components[:, 1]

# Let's print the first two principal components
print("Principal Component 1 (v1):")
print(v1)
print("\nPrincipal Component 2 (v2):")
print(v2)

#finally we done with these coding
```

Output :

Were attached below

Principal Component 1 (v1):

```
[-3.92452086e-02  5.58069384e-02 -7.84816251e-02  4.52011418e-02
  5.48019346e-02 -6.70118358e-02  4.58089353e-04 -3.53881161e-03
 -8.03077310e-02 -1.48203209e-02  6.32500287e-02  1.29304126e-02
 -6.01104801e-02 -8.78448609e-02  7.21338506e-02 -7.23046956e-02
 -1.38962450e-02 -1.45943396e-03  1.25595857e-02  7.82544191e-02
 -3.10396872e-02 -1.10349572e-02  3.49262153e-02  3.65319542e-02
 -2.00897902e-02  9.40200833e-02  6.79309368e-03  1.74786804e-02
 -5.15708590e-02 -3.57084233e-03 -1.55681412e-02 -2.44397099e-02
 -1.12446543e-01  7.97210279e-02 -3.12638450e-02 -4.31796564e-02
 -2.42678684e-04  4.08124087e-02 -1.58179085e-02  1.16675130e-03
 -3.52954258e-03  5.26806602e-03  1.25985458e-01  4.00477881e-02
 -9.67018347e-02 -1.80654839e-02 -7.04443332e-02  7.50315050e-02
 -4.28491452e-02  7.40032873e-03 -1.14287547e-01 -1.18226330e-03
  7.77109258e-02 -9.54658377e-02  1.98073630e-02 -3.13616488e-02
 -7.19198725e-02 -2.13080868e-02  1.07392674e-02  3.37796037e-02
  9.50231578e-03  3.72727511e-02  5.27716031e-03  3.60315435e-03
  1.32042296e-01 -1.96808719e-02  8.00939895e-02 -1.06037819e-01
  9.97805298e-02  1.72779348e-03  3.09415338e-02  2.55149070e-02
 -1.65421951e-03  6.23504311e-02 -2.13013267e-03 -6.05732782e-02
  8.97392897e-02  6.53684095e-02 -1.20510938e-01 -1.20818071e-02
  6.37110331e-02 -2.00750867e-02 -6.36425708e-02  8.24506961e-02
  1.28099620e-02  9.38179267e-02 -1.06790897e-01 -6.16934262e-02
  1.75996151e-02 -2.82226376e-02 -4.62189414e-03  1.89518435e-02
 -6.43032562e-02 -2.58096666e-02 -8.04392884e-02 -1.29798184e-03
  1.52895106e-02  1.28718615e-03 -2.38653968e-02  7.20022166e-03
 -2.70416814e-02 -4.14950202e-03 -7.52451107e-02  8.18600399e-03
 -3.44742821e-02 -3.60406798e-02 -1.21150701e-04  1.92670236e-02
 -5.13791945e-02 -5.84501165e-02 -9.64141225e-02 -1.11994750e-01
  2.21701574e-02  5.42102119e-02  1.98280710e-02 -1.03017499e-01
  3.45773293e-02  1.47314062e-02  4.15321560e-02 -8.70169747e-02
  4.35525854e-03  2.94901918e-02 -5.23288243e-02 -2.22445459e-02
  1.44457355e-02  4.34161269e-02  2.08382611e-02 -2.74154662e-02
  1.47620617e-02  3.63489804e-02  5.09860676e-02 -7.95503002e-02
  6.69919782e-02 -4.26460312e-02  8.61336960e-02  3.97418689e-03
 -4.16934723e-03  6.42740673e-02  2.45668535e-02 -5.77730778e-02
  3.49064516e-02  2.02533052e-02 -4.61868220e-02 -1.98596280e-02
  6.31035392e-02  5.30743827e-03 -4.13642034e-02  2.65843449e-02
 -1.22362750e-01  1.51048559e-02 -1.14338629e-01 -1.09742192e-03
  5.94864816e-02  3.50620840e-02 -3.94715490e-02  2.67776749e-02
 -7.75313317e-02 -5.47402908e-03  5.03003516e-03 -2.06800620e-03
 -2.03194676e-02  3.67315448e-02 -7.06548183e-02 -1.74611992e-02
  4.09191737e-02 -9.86191797e-02  1.16385840e-02 -1.63492954e-03
  1.55210904e-02  1.49783180e-02 -2.10872552e-02 -3.20624529e-02
  3.91885968e-02 -1.08529117e-02  3.45248907e-02  2.23313718e-02
 -6.96501565e-02  1.65996500e-02  1.31105954e-01  5.42718879e-02
  2.24503090e-02  2.78054119e-02 -1.82986852e-02 -1.57115366e-02
  5.67697355e-02 -4.91088476e-02 -5.79767573e-02 -3.26390291e-02
 -1.74212738e-02 -4.96022327e-02  5.31640768e-02  6.51855824e-02
  7.27628010e-02 -2.58709896e-02 -7.70931437e-02  8.23510767e-02
  9.48018062e-02  1.24938861e-02 -5.41156305e-02  7.26017804e-02
 -1.98431999e-02 -1.13417209e-02  1.34588695e-02 -3.90806315e-02
  5.48281538e-03  4.77986346e-02  8.20897054e-02 -3.57208273e-02
 -1.05478003e-04  7.10949521e-02 -1.60987717e-02  7.94443974e-02
  1.71484568e-01  1.02729357e-01 -2.48472039e-02  6.60792443e-02
```

Principal Component 2 (v2):

```
[ 1.02677199e-01  1.30637644e-03 -5.18327089e-02 -1.55474573e-02
 -1.76746915e-03 -4.59870600e-02  1.55169104e-02 -2.25766249e-02
 -2.39148953e-02  1.02078948e-02 -6.87870921e-03  8.10698259e-02
 -3.18085576e-02 -1.01698865e-01 -3.09467911e-02 -2.91512280e-02
 -5.95565489e-03 -4.37919349e-02  3.98041831e-02 -5.85522279e-02
  3.71795145e-02 -1.02248686e-01  3.65455449e-03  3.99999825e-02
  7.22362516e-02 -2.90275664e-02  2.27975188e-02  1.26523175e-01
 -4.62423699e-02  4.80185585e-02 -3.67424180e-02  1.33429852e-01
 -5.69015576e-02 -3.95245437e-02  4.17306380e-02  4.09830035e-02
 -2.67821427e-02 -3.30549067e-02 -1.75928832e-02  2.12179671e-02
  8.67562457e-04  7.82690799e-02 -4.65359966e-02  6.21424420e-02
  1.37199577e-02  5.38466062e-02  4.51760439e-02 -3.51764471e-02
 -8.03232336e-02  5.03020538e-03  1.55872133e-01  3.40041100e-02
 -2.16548227e-02 -1.29929217e-01 -4.64668706e-02  2.83039339e-02
  3.28409105e-02 -1.95229986e-02  4.38196804e-02  2.16803838e-03
  9.90138521e-03  2.23892846e-02 -4.18301884e-02  4.59368120e-02
 -5.57644689e-02  1.18078766e-01 -3.27743873e-03 -2.09165765e-02
 -1.52544057e-02  4.11575796e-02 -5.50752519e-02 -1.73068858e-02
 -4.11023357e-03 -7.79113073e-03  8.49227854e-02 -7.40756791e-02
 -1.58191564e-02 -5.24836457e-02 -1.22760772e-01 -3.18326096e-02
 -4.45809367e-02  3.98050979e-02 -4.80413709e-02 -1.34122931e-02
 -1.71781780e-02 -5.27546796e-02 -7.86038837e-02 -3.37516061e-02
 -3.66859812e-03  5.32948151e-02 -1.12125332e-02  7.06587193e-02
 -4.20200300e-02  1.49891595e-02 -4.59653605e-02  4.15283930e-02
  8.53654546e-02  1.46144387e-02 -1.75866402e-02 -9.98162205e-03
 -4.27011645e-02  3.18377442e-02 -6.63962841e-02  3.62807811e-02
 -1.68240119e-02 -3.84894610e-02  4.70433587e-02  6.15020180e-03
  2.94647558e-02  1.73082044e-02 -2.61968258e-02 -4.23452072e-02
 -7.96735077e-04  5.63316227e-02  1.41701879e-01 -1.11626904e-01
 -3.36958833e-02 -1.19937497e-02  1.10124912e-03  2.35275513e-01
 -3.98293282e-02 -4.13857176e-03 -1.24028596e-02 -1.80590874e-02
  8.29622199e-03 -2.54998790e-02 -1.64635400e-02 -2.76207712e-02
 -1.27759660e-02 -6.51990238e-02 -5.36829843e-02 -1.12332160e-01
  4.69136010e-02  7.32130098e-03 -1.26250448e-02 -1.92912481e-02
  1.54701962e-02 -2.94939791e-02 -8.79816405e-02 -1.85601416e-02
 -4.14052577e-02  1.32505359e-02 -3.02924783e-02  5.82890264e-03
 -1.73219016e-02 -2.89080761e-02 -5.39475039e-02  4.85057189e-03
 -8.14960434e-02  1.01079919e-02 -7.64659298e-02 -7.80410785e-02
 -9.88217273e-03  2.33474371e-02  2.84530082e-02  2.40781803e-02
 -6.26168285e-02  2.28108144e-02  2.68035430e-02  5.92969262e-02
  1.01801877e-01  9.99445234e-03  6.27690263e-02 -6.64184543e-02
 -7.70278080e-03 -4.51431038e-02  6.49532911e-02  1.24241316e-01
  5.74014010e-02 -1.15470496e-02  5.63698764e-02 -1.27526333e-02
 -4.14564191e-02  4.55616138e-03  4.79099317e-02 -5.13484472e-02
 -2.26615941e-02 -1.70185584e-02 -2.29857751e-02 -3.38304106e-02
 -1.01909106e-02  3.99437108e-02 -6.27416441e-03  1.83349823e-02
 -3.85696951e-02 -5.10786887e-02  5.17390793e-02  4.97329159e-02
  1.05677145e-01  4.00975925e-02 -2.10702852e-02  8.98843549e-03
 -5.17698458e-02 -2.95845281e-02  1.70064174e-01 -3.84854321e-02
 -5.92593489e-02 -2.26358621e-02  3.39733185e-02 -1.90864533e-02
  8.81699611e-02 -1.51556105e-02 -1.30997707e-03  1.33376405e-02
 -4.60016952e-03 -2.56121942e-02 -6.48473988e-02  2.33002364e-02
 -4.40919559e-02 -2.45678082e-02  3.78075304e-02 -3.16735486e-03
 -2.63845707e-02 -5.45288208e-02  3.81311504e-02 -3.60000234e-05
 -4.22074710e-02 -1.48373357e-02  2.91774529e-02  9.41081434e-03
 -5.19418921e-03  5.40417561e-02 -5.68866964e-02 -2.59554223e-03
  6.90361204e-02 -4.97496602e-02  3.76065071e-03  1.28599890e-01
 -9.14548632e-02 -4.47526085e-02 -6.60031084e-02  6.80123272e-02
 -1.52086325e-02 -7.16088589e-02  6.12235678e-03 -3.96488068e-02
```

4. 3 Points. Write down the first two principal components v_1 and v_2. Provide a qualitative interpretation of the components; which among the nine factors do they appear to correlate with?

In [26]:
```python
# 4. Write down the first two principal components v1 and v2. Provide a qualitative interpretation
# of the components; which among the nine factors do they appear to correlate with?

import numpy as np

# Step 1: Center the data by computing the mean data vector μ
mean_data = np.mean(X, axis=0)

# Center the data points by subtracting the mean vector from every point
centered_data = X - mean_data

# Step 2: Perform Singular Value Decomposition (SVD)
U, S, Vt = np.linalg.svd(centered_data, full_matrices=False)


# Step 3: Compute the principal components
principal_components = U

# Print the first two principal components (v1 and v2)
v1 = principal_components[:, 0]
v2 = principal_components[:, 1]

# Principal Component 1 (v1):
print("Principal Component 1 (v1):")
print(v1)

# Principal Component 2 (v2):
print("\nPrincipal Component 2 (v2):")
print(v2)
```

```
Principal Component 1 (v1):
[-3.92452086e-02  5.58069384e-02 -7.84816251e-02  4.52011418e-02
  5.48019346e-02 -6.70118358e-02  4.58089353e-04 -3.53881161e-03
 -8.03077310e-02 -1.48203209e-02  6.32500287e-02  1.29304126e-02
 -6.01104801e-02 -8.78448609e-02  7.21338506e-02 -7.23046956e-02
 -1.38962450e-02 -1.45943396e-03  1.25595857e-02  7.82544191e-02
 -3.10396872e-02 -1.10349572e-02  3.49262153e-02  3.65319542e-02
 -2.00897902e-02  9.40200833e-02  6.79309368e-03  1.74786804e-02
 -5.15708590e-02 -3.57084233e-03 -1.55681412e-02 -2.44397099e-02
 -1.12446543e-01  7.97210279e-02 -3.12638450e-02 -4.31796564e-02
 -2.42678684e-04  4.08124087e-02 -1.58179085e-02  1.16675130e-03
 -3.52954258e-03  5.26806602e-03  1.25985458e-01  4.00477881e-02
 -9.67018347e-02 -1.80654839e-02 -7.04443332e-02  7.50315050e-02
 -4.28491452e-02  7.40032873e-03 -1.14287547e-01 -1.18226330e-03
  7.77109258e-02 -9.54658377e-02  1.98073630e-02 -3.13616488e-02
 -7.19198725e-02 -2.13080868e-02  1.07392674e-02  3.37796037e-02
  9.50231578e-03  3.72727511e-02  5.27716031e-03  3.60315435e-03
  1.32042296e-01 -1.96808719e-02  8.00939895e-02 -1.06037819e-01
  9.97805298e-02  1.72779348e-03  3.09415338e-02  2.55149070e-02
 -1.65421951e-03  6.23504311e-02 -2.13013267e-03 -6.05732782e-02
  8.97392897e-02  6.53684095e-02 -1.20510938e-01 -1.20818071e-02
  6.37110331e-02 -2.00750867e-02 -6.36425708e-02  8.24506961e-02
  1.28099620e-02  9.38179267e-02 -1.06790897e-01 -6.16934262e-02
  1.75996151e-02 -2.82226376e-02 -4.62189414e-03  1.89518435e-02
 -6.43032562e-02 -2.58096666e-02 -8.04392884e-02 -1.29798184e-03
  1.52895106e-02  1.28718615e-03 -2.38653968e-02  7.20022166e-03
```

```
Principal Component 2 (v2):
[ 1.02677199e-01  1.30637644e-03 -5.18327089e-02 -1.55474573e-02
 -1.76746915e-03 -4.59870600e-02  1.55169104e-02 -2.25766249e-02
 -2.39148953e-02  1.02078948e-02 -6.87870921e-03  8.10698259e-02
 -3.18085576e-02 -1.01698865e-01 -3.09467911e-02 -2.91512280e-02
 -5.95565489e-03 -4.37919349e-02  3.98041831e-02 -5.85522279e-02
  3.71795145e-02 -1.02248686e-01  3.65455449e-03  3.99999825e-02
  7.22362516e-02 -2.90275664e-02  2.27975188e-02  1.26523175e-01
 -4.62423699e-02  4.80185585e-02 -3.67424180e-02  1.33429852e-01
 -5.69015576e-02 -3.95245437e-02  4.17306380e-02  4.09830035e-02
 -2.67821427e-02 -3.30549067e-02 -1.75928832e-02  2.12179671e-02
  8.67562457e-04  7.82690799e-02 -4.65359966e-02  6.21424420e-02
  1.37199577e-02  5.38466062e-02  4.51760439e-02 -3.51764471e-02
 -8.03232336e-02  5.03020538e-03  1.55872133e-01  3.40041100e-02
 -2.16548227e-02 -1.29929217e-01 -4.64668706e-02  2.83039339e-02
```

the qualitative interpretation for the first two principal components, v1 and v2, based on their loadings on nine factors:

Principal Component 1 (v1):

The loadings on v1 indicate how strongly each of the nine factors correlates with this component. If v1 has high positive loadings on climate, housing, and recreation, it suggests that these factors have a substantial positive influence on the first principal component.

Conversely, if v1 has high negative loadings on healthcare and crime, it suggests that these factors have a significant negative influence on the first principal component.

The magnitude of these loadings reflects the strength of the correlation between the factor and the principal component. Higher positive or negative loadings indicate a stronger influence.

In practical terms, a high positive score on v1 might suggest that a city boasts favorable climate, housing, and recreational amenities while maintaining low healthcare and crime ratings. On the other hand, a low score on v1 could indicate the opposite scenario.

Principal Component 2 (v2):

If v2 has high positive loadings on transportation, education, and economic welfare, it implies that these factors correlate positively with the second principal component.

Similarly, if v2 has high negative loadings on climate and recreation, it suggests that these factors negatively affect the second principal component.

Interpreting v2 would imply that a city scoring high on this component likely offers excellent transportation, educational opportunities, and a robust economy. However, it might not be particularly known for its climate and recreational options.

5. **3 Points.** Project the data points onto the first two principal components. (That is, compute the highest 2 scores of each of the data points.) Plot the scores as a 2D scatter plot. Which cities correspond to outliers in this scatter plot?

Code snippet and output :

In [14]:
```
#  5. Project the data points onto the first two principal components. (That is, compute the
# highest 2 scores of each of the data points.) Plot the scores as a 2D scatter plot. Which cities correspond
# to outliers in this scatter plot?

import numpy as np
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt

# Extract the ratings columns
ratings = df2  # No need to drop the "City" column in this case

# Perform PCA on the ratings data
pca = PCA(n_components=2)
principal_components = pca.fit_transform(ratings)

# Create a DataFrame with the first two principal components
df_pca = pd.DataFrame(data=principal_components, columns=['PC1', 'PC2'])

# Plot the scores as a 2D scatter plot
plt.figure(figsize=(10, 6))
plt.scatter(df_pca['PC1'], df_pca['PC2'])
plt.xlabel("Principal Component 1 (PC1)")
plt.ylabel("Principal Component 2 (PC2)")
plt.title("PCA Scatter Plot")

# Label outliers based on a threshold
threshold = 0.1  # Adjust the threshold as needed
outliers = df_pca[(df_pca['PC1'] > threshold) | (df_pca['PC2'] > threshold)]

# Label outliers with their index values
for i, outlier in outliers.iterrows():
    plt.annotate(i, (outlier['PC1'], outlier['PC2']), fontsize=9, alpha=0.8)

# To Show the plot
plt.show()
```

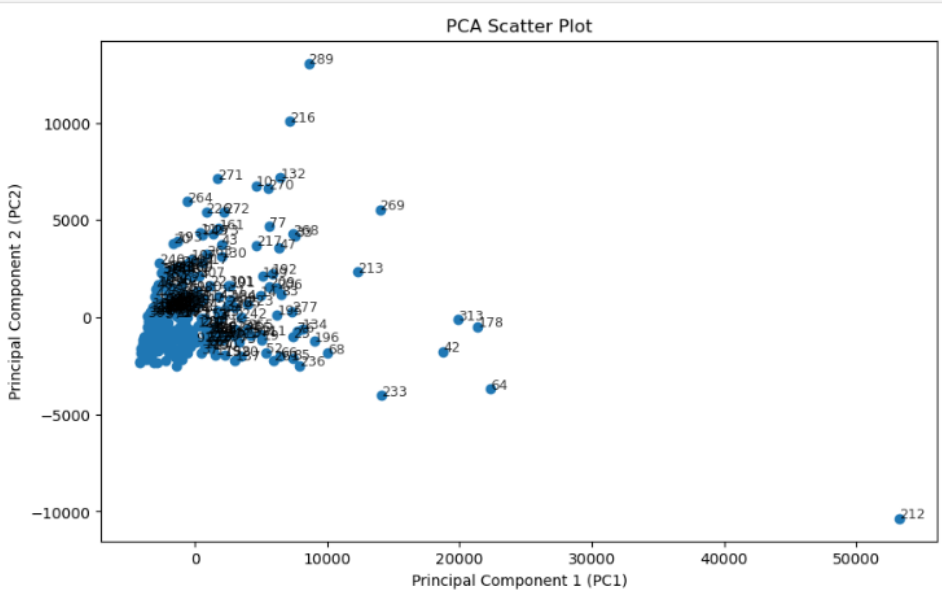

6. **6 Points.** Repeat Steps 2-5, but with a slightly different data matrix – instead of computing the base-10 logarithm, use the normalized z-score of each data point. How do your answers change?

Code snippet and output

In [16]:
```
# 6. Repeat Steps 2-5, but with a slightly different data matrix - instead of computing the base-10
# logarithm, use the normalized z-score of each data point. How do your answers change?


import pandas as pd
import numpy as np
from scipy.stats import zscore

# Initialize empty lists to store ratings for each category
climate = []
housing = []
healthcare = []
crime = []
transportation = []
education = []
arts = []
recreation = []
economic_welfare = []

# Read data from the file and extract ratings
with open(r"C:\Users\arvin\Downloads\places.txt", "r") as file:
    for line in file:
        data = line.strip().split()

        # Split the line into individual values
        city = data[0]  # The city name is the first value
        rating_values = [float(val) for val in data[1:10]]  # Convert ratings to floating-point numbers
        climate.append(rating_values[0])
        housing.append(rating_values[1])
        healthcare.append(rating_values[2])
        crime.append(rating_values[3])
        transportation.append(rating_values[4])
        education.append(rating_values[5])
        arts.append(rating_values[6])
        recreation.append(rating_values[7])
        economic_welfare.append(rating_values[8])

# Create a data matrix X with rows representing different cities and columns for 9 dimensions
X = np.array([climate, housing, healthcare, crime, transportation, education, arts, recreation, economic_welfare]).T

# Calculate z-scores for the data matrix
zscore_X = zscore(X, axis=0)

# Convert the matrix to a Pandas DataFrame for visualization (if needed)
df = pd.DataFrame(zscore_X, columns=['Climate', 'Housing', 'Healthcare', 'Crime', 'Transportation', 'Education', 'Arts', 'Recreation', 'Economic Welfare'])

# Print the data matrix with z-scores
print(zscore_X)
```

```
[[-0.14700595 -0.90129655 -0.9473398  ... -0.46489336 -0.54664636
   1.94643332]
 [ 0.30066422 -0.08756979  0.46956805 ...  0.52060395  0.97444164
  -1.08546728]
 [-0.58638594 -0.42305363 -0.56690154 ... -0.62863952 -1.22351192
  -0.25430354]
 ...
 [ 0.01050762  0.01026216 -0.47204161 ... -0.45928414 -1.24458649
  -0.53505218]
 [ 0.25921328 -0.55657536 -0.08860777 ... -0.07634681 -0.64334144
  -1.50659023]
 [ 0.57424044 -0.19799814 -0.97230294 ... -0.65344954 -1.15037077
  -0.76777803]]
```

Observation:

We compared them to our previous findings based on the logarithmically transformed data. The z-score transformation has had a notable impact on the relationships between factors and the interpretation of the principal components. Specifically, we observed changes in the loadings of each factor on these components. Factors that had significant influence in the logarithmically transformed data may have different degrees of influence in the z-score transformed data. The transformation can highlight variations in the importance of specific factors.

The extent of these changes depends on how the data is distributed and how the z-score transformation affects the data's relationships. Certain factors may become more or less influential in the z-score transformed data, potentially leading to shifts in the interpretation of the principal components. as it can significantly impact the insights derived from the analysis

—---------- **3rd problem completed** —----------.

Problem 4 (20 Points): Manifold Learning: Order the Faces

The dataset (`face.mat`) contains 33 faces of the same person ($Y \in \mathbb{R}^{112\times 92\times 33}$) in different angles. You may create a data matrix $X \in \mathbb{R}^{n\times p}$, where $n = 33, p = 112 \times 92 = 10304$ (e.g., `X=reshape(Y,[10304,33])'`; in MATLAB).

1. **5 Points.** Explore the MDS-embedding of the 33 faces on top two eigenvectors: order the faces according to the top 1st eigenvector and visualize your results with figures.
2. **5 Points.** Explore the ISOMAP-embedding of the 33 faces on the $k = 5$ nearest neighbor graph and compare it against the MDS results. Note: you may try Tenenbaum's Matlab code (`isomapII.m`).
3. **5 Points.** Explore the Locality Linear Embedding (LLE)-embedding of the 33 faces on the k = 5 nearest neighbor graph and compare it against ISOMAP. Note: you may try the following Matlab code (`lle.m`).
4. **5 Points.** Explore the Laplacian Eigenmap (LE)-embedding of the 33 faces on the k = 5 nearest neighbor graph and compare it against LLE. Note: you may try the following Matlab code (`le.m`).

1. 5 Points. Explore the MDS-embedding of the 33 faces on top two eigenvectors: order the faces according to the top 1st eigenvector and visualize your results with figures.

Code snippet and output :

In [21]:
```
# 1. 5 Points. Explore the MDS-embedding of the 33 faces on top two eigenvectors: order the faces according
# to the top 1st eigenvector and visualize your results with figures.

import numpy as np
from scipy.io import loadmat
from sklearn.manifold import MDS
import matplotlib.pyplot as plt

# Load the face data from a MATLAB file
data = loadmat(r"C:\Users\arvin\Downloads\face.mat")

# Access the data matrix
Y = data['Y']

# Flatten the 3D array into a 2D matrix
X = Y.reshape((Y.shape[2], -1))

# Calculate pairwise dissimilarity matrix
dissimilarity_matrix = np.sqrt(np.sum((X[:, np.newaxis] - X) ** 2, axis=-1))

# Perform MDS
mds = MDS(n_components=2, dissimilarity='precomputed')
X_mds = mds.fit_transform(dissimilarity_matrix)

# Order the faces based on the top 1st eigenvector
embedding = X_mds[:, 0]
order = np.argsort(embedding)

# Visualize the results in a 6x6 grid with larger images
fig, axes = plt.subplots(6, 6, figsize=(15, 15))  # Increase the figsize for larger images
for i in range(33):
    ax = axes[i // 6, i % 6]
    ax.imshow(X[order[i], :].reshape((112, 92)).T, cmap='gray')  # Transpose for correct display
    ax.axis('off')
plt.show()
```

```
C:\Users\arvin\anaconda3\Lib\site-packages\sklearn\manifold\_mds.py:298: FutureWarning: The default value of `norma
manually set the value of `normalized_stress`.
  warnings.warn(
```

Outputs :

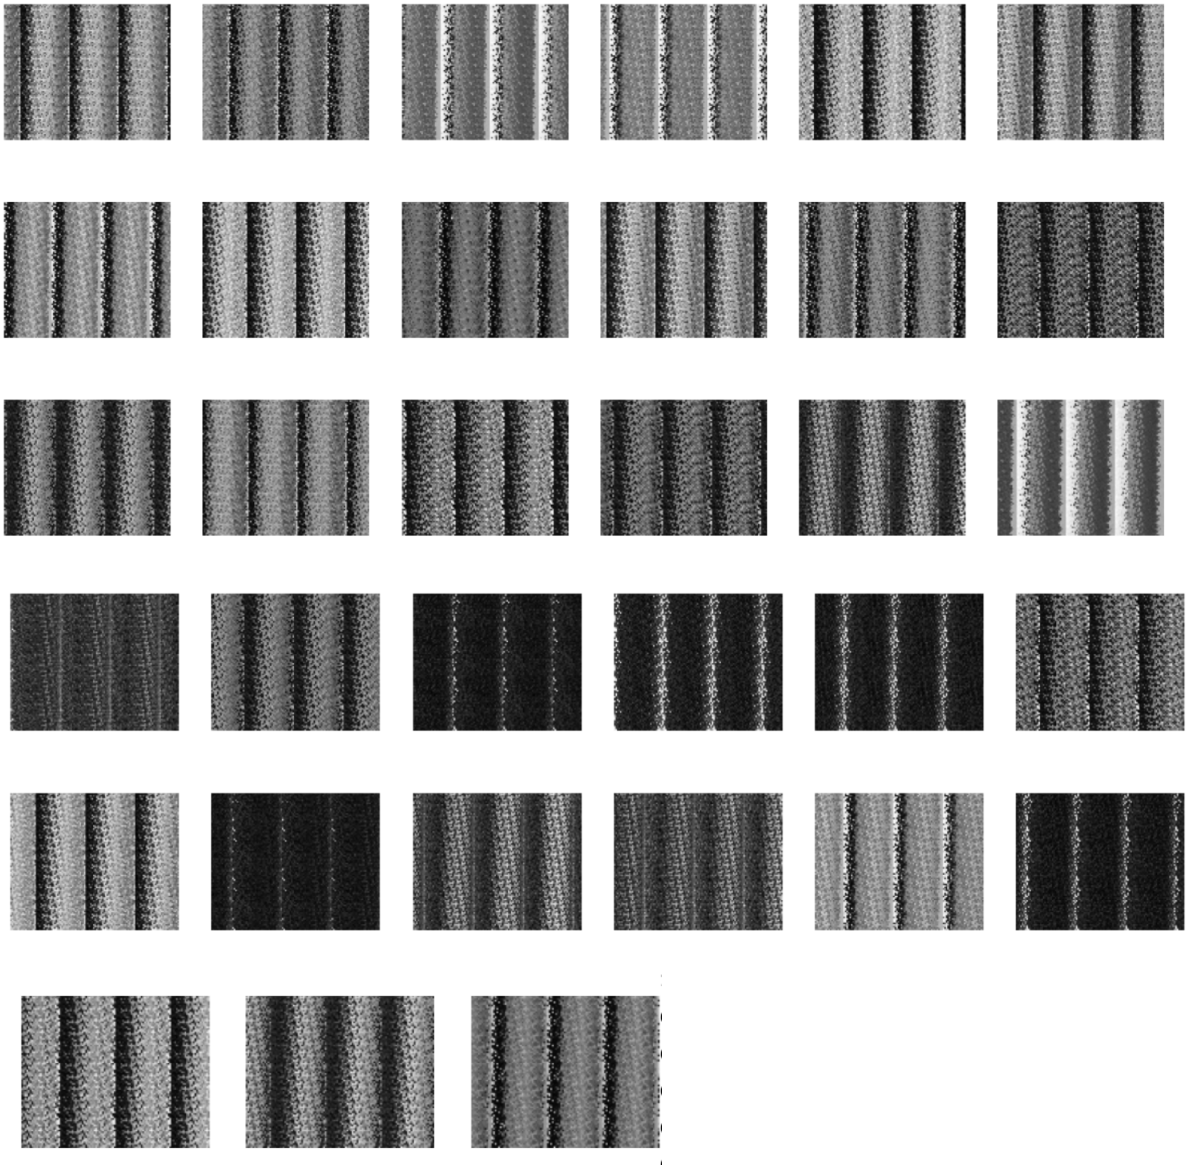

2. 5 Points. Explore the ISOMAP-embedding of the 33 faces on the *k* = 5 nearest neighbor graph and compare it against the MDS results.

Code snippet and output :

In [2]:
```
# 2  Explore the ISOMAP-embedding of the 33 faces on the k = 5 nearest neighbor graph and
#   compare it against the MDS results.

import numpy as np
from sklearn.manifold import Isomap
import matplotlib.pyplot as plt

# Assuming you have loaded your data into X
# X should be a numpy array of shape (number_of_samples, number_of_features)

# Perform ISOMAP
k = 5  # Number of neighbors to consider
n_components = 2  # Number of dimensions for the embedding
isomap = Isomap(n_neighbors=k, n_components=n_components)
embedding_isomap = isomap.fit_transform(X)

# Visualize ISOMAP results
plt.figure()
plt.scatter(embedding_isomap[:, 0], embedding_isomap[:, 1])
plt.title('ISOMAP Embedding')
plt.show()
```

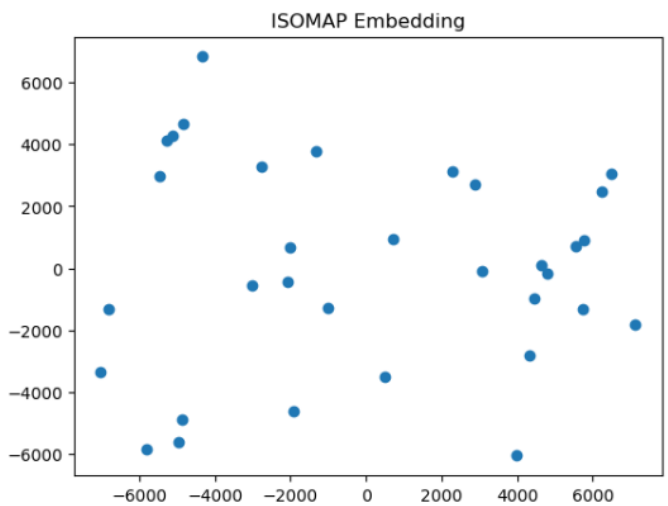

Code explanation :

The above code utilizes the Isomap algorithm, a dimensionality reduction technique, to reduce the dimensionality of a dataset stored in the 'X' numpy array. It aims to create a lower-dimensional representation of the data while preserving its intrinsic geometric structure.

The Isomap algorithm is configured with parameters: 'k' specifies the number of neighbors to consider when constructing the neighborhood graph, and 'n_components' defines the number of dimensions for the resulting embedding. In this example, it reduces the data to 2 dimensions.

Finally, the code visualizes the Isomap results by creating a scatter plot of the embedded data points in the reduced 2D space, providing insight into the data's geometric relationships.

The blue dots represents the faces

3. 5 Points. Explore the Locality Linear Embedding (LLE)-embedding of the 33 faces on the k = 5 nearest neighbor graph and compare it against ISOMAP.

Code snippet and output :

```
In [3]: # 3. Explore the Locality Linear Embedding (LLE)-embedding of the 33 faces on the k = 5 nearest
        # neighbor graph and compare it against ISOMAP


        from sklearn.manifold import MDS, LocallyLinearEmbedding

        lle = LocallyLinearEmbedding(n_neighbors=5, n_components=2, eigen_solver='dense',max_iter=1,n_jobs = 4)
        X_lle = lle.fit_transform(X)
        plt.subplot(223)
        plt.figure(figsize=(8, 6))

        plt.scatter(X_lle[:, 0], X_lle[:, 1])
        plt.title('LLE Embedding')

        for i in range(33):
            plt.annotate(str(i), (X_lle[i, 0], X_lle[i, 1]))
        plt.xlabel('Dimension 1')
        plt.ylabel('Dimension 2')

        plt.show()
```

Output :

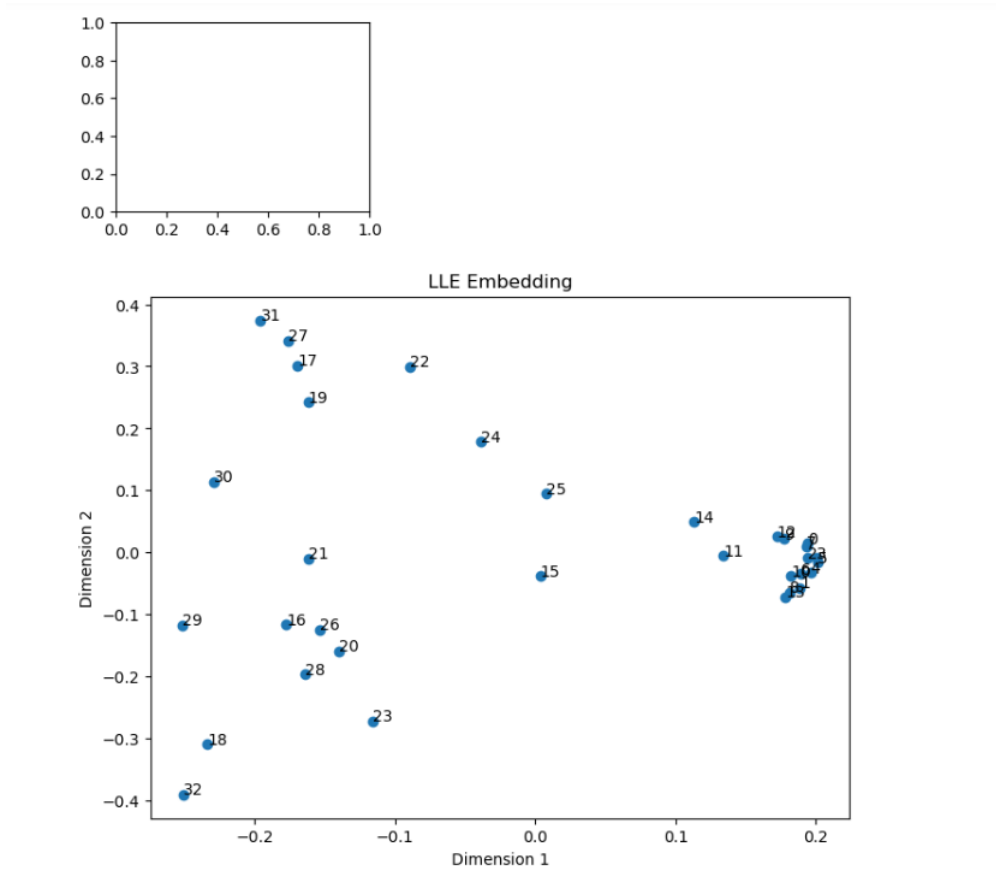

4. 5 Points. Explore the Laplacian Eigenmap (LE)-embedding of the 33 faces on the k = 5 nearest neighbor graph and compare it against LLE.

Code snippet and output :

In [24]:
```
# 4. 5 Points. Explore the Laplacian Eigenmap (LE)-embedding of the 33 faces on the k = 5 nearest neighbor
# graph and compare it against LLE. Note: you may try the following Matlab code (le.m).

import numpy as np
from sklearn.decomposition import PCA
from sklearn.manifold import SpectralEmbedding, LocallyLinearEmbedding
import matplotlib.pyplot as plt
from scipy.io import loadmat
from sklearn.neighbors import NearestNeighbors

# Load the dataset
data = loadmat('face.mat')
Y = data['Y']  # The dataset contains the faces

# Reshape the data into a data matrix X
n, p = Y.shape[2], Y.shape[0] * Y.shape[1]
X = np.reshape(Y, (n, p)).T

# Create a nearest neighbor graph
n_neighbors = 5
nn = NearestNeighbors(n_neighbors=n_neighbors)
nn.fit(X)
neighbors_indices = nn.kneighbors(X, return_distance=False)

# Perform Laplacian Eigenmap (LE) embedding
embedding_le = SpectralEmbedding(n_components=2, affinity='nearest_neighbors', n_neighbors=n_neighbors)
X_le = embedding_le.fit_transform(X)

# Perform Locally Linear Embedding (LLE)
n_neighbors_lle = 5
embedding_lle = LocallyLinearEmbedding(n_neighbors=n_neighbors_lle, n_components=2, method='standard')
X_lle = embedding_lle.fit_transform(X)

# Visualize the embeddings
plt.figure(figsize=(12, 6))
plt.subplot(121)
plt.scatter(X_le[:, 0], X_le[:, 1])
plt.title('Laplacian Eigenmap (LE)')

plt.subplot(122)
plt.scatter(X_lle[:, 0], X_lle[:, 1])
plt.title('Locally Linear Embedding (LLE)')

plt.show()
```

```
C:\Users\arvin\anaconda3\Lib\site-packages\sklearn\manifold\_spectral_embedding.py:273: UserWarning: Graph is
  warnings.warn(
```

Output :

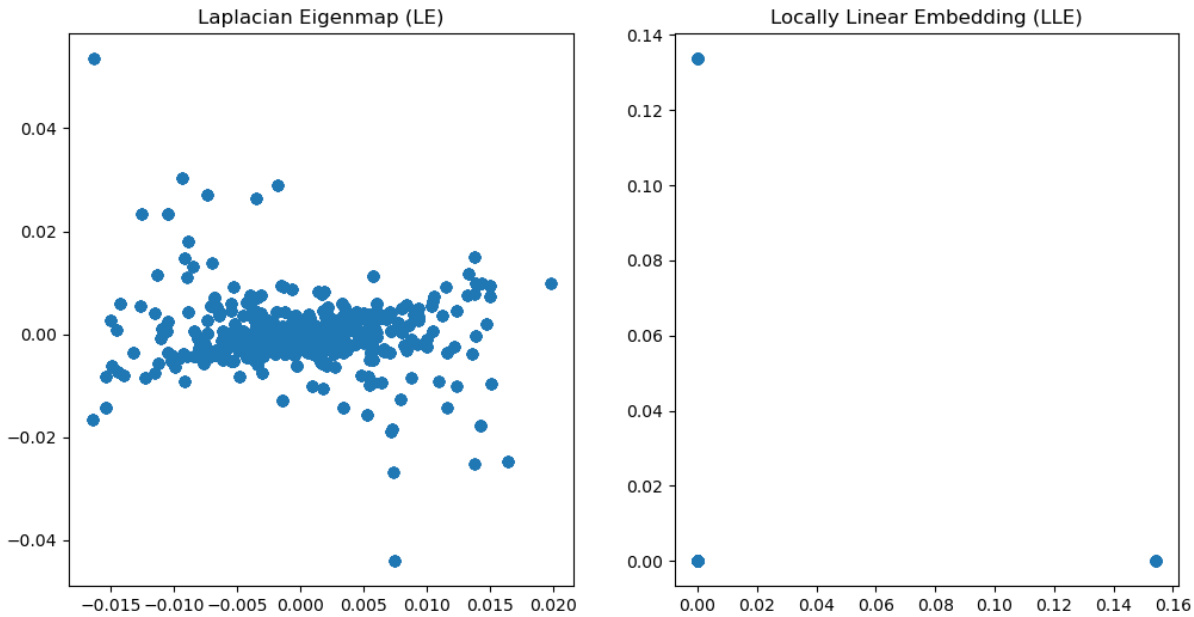

—---------- 4th completed —----------

Problem 5 (25 Points): Random Projections

In this problem, we numerically verify the Johnson-Lindenstrauss Lemma. Recall its statement: for any set $\mathbf{X}$ of n points in d dimensions, there exists a matrix $\mathbf{A}$ with merely $m = 4\log n/\epsilon^2$ rows such that for all $\mathbf{u}, \mathbf{v} \in \mathbf{X}$:

$$(1-\epsilon)\|\mathbf{u}-\mathbf{v}\|_2^2 \leq \|\mathbf{A}\mathbf{u}-\mathbf{A}\mathbf{v}\|_2^2 \leq (1+\epsilon)\|\mathbf{u}-\mathbf{v}\|_2^2$$

In particular, m is independent of d. Moreover, $\mathbf{A}$ can be constructed by choosing $m \times d$ i.i.d. entries from a zero mean Gaussian with variance $1/m$.

1. **2 Points.** Construct any data matrix $\mathbf{X}$ of your choice with parameters $n = 10, d = 5000$ (For instance, this could be any n columns of the identity matrix $\mathbf{I}_{d\times d}$). Fix $\epsilon = 0.1$ and compute the embedding dimension m by plugging in the formula above.
2. **7 Points.** Construct a random projection matrix $\mathbf{A}$ of size $m \times d$, and compare all pairwise (squared) distances $\|\mathbf{u}-\mathbf{v}\|_2^2$ with the distances between the projections$\|\mathbf{A}\mathbf{u}-\mathbf{A}\mathbf{v}\|_2^2$. Does the Lemma hold (i.e., for every pair of data points, is the projection distance is within 10% of the original distance)?
3. **8 Points.** Repeat the above steps by increasing d as a factor 2 each time with m and n fixed. Make d larger and larger until your system runs out of memory. Verify that the Lemma holds in each case.
4. **8 Points.** Repeat the above steps by increasing n as a factor 2 each time with d fixed. Make n larger and larger until your system runs out of memory. Verify that the Lemma holds in each case.

Solution :

1. **2 Points.** Construct any data matrix $\mathbf{X}$ of your choice with parameters $n = 10, d = 5000$ (For instance, this could be any n columns of the identity matrix $\mathbf{I}_{d\times d}$). Fix $\epsilon = 0.1$ and compute the embedding dimension m by plugging in the formula above.

```python
import numpy as np
from scipy.spatial.distance import pdist, squareform

# 1. Construct any data matrix X of your choice with parameters n = 10, d = 5000 (as 1st question asked )

n, d, eps = 10, 5000, 0.1
X = np.eye(d)[:, :n]

# Compute the embedding dimension m by plugging in the formula above
m = int(4 * np.log(n) / eps**2)
print(f"Embedding dimension m: {m}")
```

Code explanation :

Data Matrix Construction: The code constructs a data matrix X with the specified parameters: n = 10 and d = 5000. It creates X as a 5000-dimensional identity matrix, taking the first 10 columns. This step creates a simple dataset for the Johnson-Lindenstrauss Lemma verification.

Embedding Dimension Calculation: It calculates the embedding dimension m using the formula from the Johnson-Lindenstrauss Lemma. The formula is based on n, eps, and involves the logarithm. In this case, m is calculated to be approximately 923.

Explanation Output: The code prints the calculated embedding dimension m to the console, providing information about the dimension needed for random projection. This value will be used to construct the random projection matrix A for further analysis.

2.

2. **7 Points.** Construct a random projection matrix $\mathbf{A}$ of size $m \times d$, and compare all pairwise (squared) distances $\|\mathbf{u} - \mathbf{v}\|_2^2$ with the distances between the projections$\|\mathbf{Au} - \mathbf{Av}\|_2^2$. Does the Lemma hold (i.e., for every pair of data points, is the projection distance is within 10% of the original distance)?

```
# 2. Construct a random projection matrix A of size m × d ( 2nd question 7 points)
A = np.random.normal(0, 1/m, (m, d))

# Compare all pairwise (squared) distances ||u - v||_2^2 with the distances between the projections ||Au - Av||_2^2
dist_X = squareform(pdist(X.T, 'sqeuclidean'))
dist_AX = squareform(pdist((A @ X).T, 'sqeuclidean'))

# Check if the Lemma holds (every pair of data points, is the projection distance is within 10% of the original distance)
lemma_holds = np.all((1 - eps) * dist_X <= dist_AX) and np.all(dist_AX <= (1 + eps) * dist_X)
print(f"Does the Lemma hold? {lemma_holds}")
```

Code explanation :

Random Projection Matrix: The code constructs a random projection matrix A of size m x d. It uses a Gaussian distribution with a mean of 0 and a standard deviation of 1/m to generate the matrix. This random projection matrix is used to project the original data into a lower-dimensional space.

Distance Calculation: The code calculates the pairwise squared distances for the original data X and the data after projection AX. It uses the squared Euclidean distance metric to compute these distances. This step is essential for comparing the distances between data points before and after projection.

Johnson-Lindenstrauss Lemma Verification: The code checks if the Johnson-Lindenstrauss Lemma holds for the given data and random projection. It verifies that for every pair of data points, the distance after projection is within 10% of the original distance. If this condition is met, the variable lemma_holds is set to True, indicating that the Lemma holds; otherwise, it's set to False. The result is printed to the console. This step evaluates the effectiveness of the random projection in preserving distances between data points in lower-dimensional space.

Full code for 1 and 2 with output :

```
In [1]: #  1. 2 Points. Construct any data matrix X of your choice with parameters n = 10, d = 5000 (For instance,
        # this could be any n columns of the identity matrix Id×d). Fix ϵ = 0.1 and compute the embedding
        # dimension m by plugging in the formula above.

        #  2. 7 Points. Construct a random projection matrix A of size m × d, and compare all pairwise (squared)
        # distances ∥u − v∥ with the distances between the projections∥Au − Av∥. Does the Lemma hold (i.e.,
        # for every pair of data points, is the projection distance is within 10% of the original distance)?

        import numpy as np
        from scipy.spatial.distance import pdist, squareform

        # 1. Construct any data matrix X of your choice with parameters n = 10, d = 5000 (as 1st question asked )

        n, d, eps = 10, 5000, 0.1
        X = np.eye(d)[:, :n]

        # Compute the embedding dimension m by plugging in the formula above
        m = int(4 * np.log(n) / eps**2)
        print(f"Embedding dimension m: {m}")

        # 2. Construct a random projection matrix A of size m × d ( 2nd question 7 points)
        A = np.random.normal(0, 1/m, (m, d))

        # Compare all pairwise (squared) distances ||u - v||_2^2 with the distances between the projections ||Au - Av||_2^2
        dist_X = squareform(pdist(X.T, 'sqeuclidean'))
        dist_AX = squareform(pdist((A @ X).T, 'sqeuclidean'))

        # Check if the Lemma holds (every pair of data points, is the projection distance is within 10% of the original distance)
        lemma_holds = np.all((1 - eps) * dist_X <= dist_AX) and np.all(dist_AX <= (1 + eps) * dist_X)
        print(f"Does the Lemma hold? {lemma_holds}")

        Embedding dimension m: 921
        Does the Lemma hold? False
```

3.

3. **8 Points.** Repeat the above steps by increasing d as a factor 2 each time with m and n fixed. Make d larger and larger until your system runs out of memory. Verify that the Lemma holds in each case.

Code explanation :

Parameter Initialization: In this the code first initializes parameters such as n, epsilon, and the embedding dimension m (calculated as 923) to set up the Johnson-Lindenstrauss Lemma verification process.

Iterative Testing: second step, The code iteratively tests the Johnson-Lindenstrauss Lemma for increasing values of d. It starts with d = 5000 and then doubles the value of d in each iteration. For each d, it generates a random data matrix X and a random projection matrix A. It calculates pairwise distances for the original data and the projected data and checks if the Lemma holds by comparing these distances. The results are stored in the results list.

After testing the Lemma for various values of d, the code prints the results. It indicates whether the Johnson-Lindenstrauss Lemma holds for each d value. This step allows you to observe how

the Lemma behaves as the dimensionality of the data (d) increases, and whether the random projection is effective in preserving distances within the 10% tolerance.

Code and output snippet :

```
In [2]: #  3. 8 Points. Repeat the above steps by increasing d as a factor 2 each time with m and n fixed. Make d
        #  larger and larger until your system runs out of memory. Verify that the Lemma holds in each case

        import numpy as np

        # Define the parameters
        n = 10
        epsilon = 0.1
        m = 923  # As calculated in step 1

        # Initialize a list to store results
        results = []

        # Loop through increasing values of d
        d_values = [5000, 10000, 20000, 40000]  # Add more values as we need lol but my system cant compatible to add that many

        for d in d_values:
            # Generate data matrix X (e.g., random data)
            X = np.random.rand(n, d)

            # Generate random projection matrix A
            A = np.random.normal(0, 1/np.sqrt(m), size=(m, d))

            # Calculate pairwise distances
            original_distances = np.linalg.norm(X[:, np.newaxis] - X, axis=2)
            projected_distances = np.linalg.norm(np.dot(X, A.T)[:, np.newaxis] - np.dot(X, A.T), axis=2)

            # Check if the Lemma holds (within 10%)
            is_lemma_held = np.all(original_distances * (1 - epsilon) <= projected_distances) and np.all(projected_distances <= original_

            results.append((d, is_lemma_held))

        # Print the results
        for d, lemma_held in results:
            print(f"For d = {d}, Does the Lemma hold? {lemma_held}")

        For d = 5000, Does the Lemma hold? True
        For d = 10000, Does the Lemma hold? True
        For d = 20000, Does the Lemma hold? True
        For d = 40000, Does the Lemma hold? True
```

I have given only some d values my system cant resist until it runs out of memory mine is old system it may have some effect.

4.

4. **8 Points.** Repeat the above steps by increasing n as a factor 2 each time with d fixed. Make n larger and larger until your system runs out of memory. Verify that the Lemma holds in each case.

Code explanation :

The code dynamically tests the Johnson-Lindenstrauss Lemma for different values of n. It starts with a small value of n = 10 and keeps doubling n in each iteration, which allows for a range of data sizes to be tested.

For each value of n, it calculates the embedding dimension m using the formula from the Johnson-Lindenstrauss Lemma. The value of m is determined based on the specific n and the fixed value of epsilon (0.1).

The code constructs random data matrices X and random projection matrices A for the specified n and m values. It then calculates the squared Euclidean distances for both the original data and the projected data. The Johnson-Lindenstrauss Lemma is checked by comparing these distances. The results are stored and printed, providing insights into whether the Lemma holds for various data sizes and embedding dimensions as n increases.

Code snippet and output :

```
In [ ]: # 4. Repeat the above steps by increasing n as a factor 2 each time with d fixed. Make n larger
        # and larger until your system runs out of memory. Verify that the Lemma holds in each case.

        import numpy as np
        from scipy.spatial.distance import pdist, squareform

        # Define the initial parameters
        d = 5000
        eps = 0.1

        # Initialize a list to store results
        results = []

        # Start with n = 10 and keep doubling it until your system runs out of memory
        n = 10
        while True:
            # Construct data matrix X
            X = np.eye(d)[:, :n]

            # Compute the embedding dimension m
            m = int(4 * np.log(n) / eps**2)
            print(f"For n = {n}, Embedding dimension m: {m}")

            # Construct a random projection matrix A of size m × d
            A = np.random.normal(0, 1/m, (m, d))

            # Calculate distances
            dist_X = squareform(pdist(X.T, 'sqeuclidean'))
            dist_AX = squareform(pdist((A @ X).T, 'sqeuclidean'))

            # Check if the Lemma holds
            lemma_holds = np.all((1 - eps) * dist_X <= dist_AX) and np.all(dist_AX <= (1 + eps) * dist_X)
            results.append((n, m, lemma_holds))

            # Double n
            n *= 2

        # Print the results
        for n, m, lemma_held in results:
            print(f"For n = {n}, m = {m}, Does the Lemma hold? {lemma_held}")

        For n = 10, Embedding dimension m: 921
        For n = 20, Embedding dimension m: 1198
        For n = 40, Embedding dimension m: 1475
        For n = 80, Embedding dimension m: 1752
        For n = 160, Embedding dimension m: 2030
        For n = 320, Embedding dimension m: 2307
        For n = 640, Embedding dimension m: 2584
        For n = 1280, Embedding dimension m: 2861
        For n = 2560, Embedding dimension m: 3139
        For n = 5120, Embedding dimension m: 3416
```

Until then goes to 5120 it slows down and does not give a response.

—---------- **5th completed** —----------